


Modifikacija i *rebuild*-ovanje Linux kernela

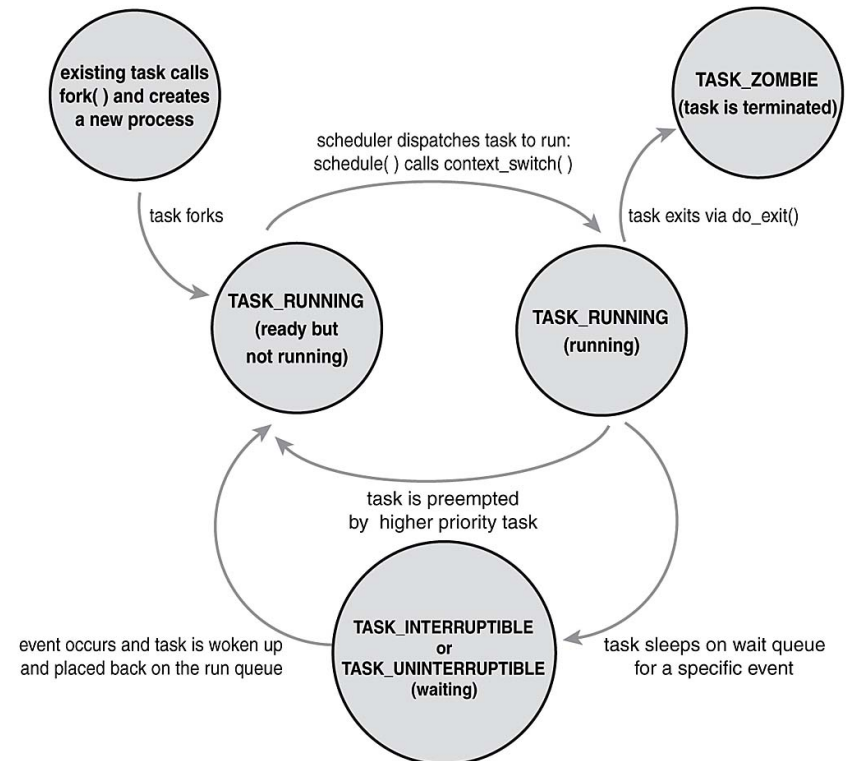
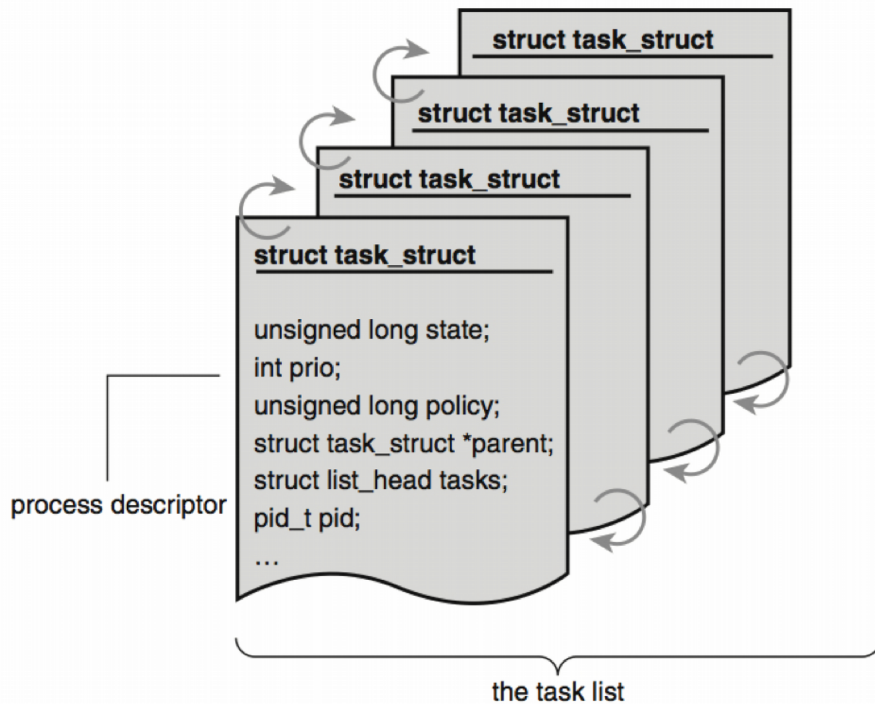
Aleksa Trajković
Filip Stamenković

370
342



Uvod

- Upravljanje procesima je jedno od glavnih zaduženja svakog operativnog sistema.
- CPU je jedan od najznačajnijih resursa u računarskom sistemu, koji konkurentni procesi moraju efikasno deliti.




Process Scheduling

- Raspoređivanje (planiranje procesa, *process scheduling*) je ključno za multiprogramiranje.
- Cilj raspoređivanja je da se procesoru dodeljuju procesi za izvršenje, na način da se zadovolje zahtevi sistema, kao što su vreme odziva i propusna moć.
- Postoji puno algoritama za raspoređivanje u literaturi, jedan od najjednostavnijih je algoritam kružnog raspoređivanja (*round-robin*).




Round-Robin (RR)



- **Funkcija selekcije:** ista kao kod FCFS
- **Mod odluke:** sa prekidanjem (*preemptive*)
 - Procesu se dopušta da se izvršava dok mu ne istekne dodeljeni vremenski period (kvant tipično 10 do 100 ms)
 - Tada se javlja prekidni signal od časovnika i proces se vraća u red spremnih procesa

Cilj domaćeg zadatka

- Dodavanje novog *scheduler*-a (“domaći”) u red već postojećih.
- Dodavanje sistemskog poziva za prikupljanje statistike o procesima koji su se izvršili na novom *scheduler*-u.



Linux Scheduler

- *Linux scheduler* je modularan i omogućava da se različitim tipovima procesa dodele različiti algoritmi raspoređivanja.
- Svaki *scheduling class* enkapsulira *scheduling policy*.
 - Real-Time Scheduler
 - Completely Fair Scheduler
 - Idle Scheduler
 - ...
- Scheduling klase su hijerarhijski organizovane.

Scheduling policy

- Prvi korak u implementaciji novog scheduler-a je definisati scheduling policy.
 - ***(kernel-source)/include/uapi/linux/sched.h***
 - ***/usr/include/bits/sched.h***

```
/*  
 * Scheduling policies  
 */  
#define SCHED_NORMAL    0  
#define SCHED_FIFO      1  
#define SCHED_RR        2  
#define SCHED_BATCH     3  
/* SCHED_ISO: reserved but not implemented yet */  
#define SCHED_IDLE      5  
#define SCHED_DEADLINE  6  
#define SCHED_DOMACI    10
```

```
#define PROCESS_BUFF_LEN 50  
  
/* Scheduling algorithms. */  
#define SCHED_OTHER      0  
#define SCHED_FIFO      1  
#define SCHED_RR        2  
#define SCHED_DOMACI    10  
#ifdef __USE_GNU  
# define SCHED_BATCH     3  
# define SCHED_IDLE      5  
  
# define SCHED_RESET_ON_FORK 0x40000000  
#endif
```

Dve najvažnije strukture, *task_struct* i *rq*

(kernel-source)/include/linux/sched.h

```
struct task_struct {  
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */  
    void *stack;  
    atomic_t usage;  
    unsigned int flags; /* per process flags, defined below */  
    unsigned int ptrace;
```

```
//domaci  
    unsigned int domaci_time_slice;  
    struct list_head domaci_list;  
    unsigned long domaci_ticks;  
    unsigned long domaci_rt;  
    //
```

(dodat je i kod za inicijalizaciju
dodatih polja strukture u
/include/linux/init_task.h i
/kernel/fork.c)

(dodat je i kod za inicijalizaciju
dodatih polja strukture u
/kernel/sched/core.c)

(kernel-source)/kernel/sched/sched.h

```
struct domaci_rq  
{  
    struct list_head domaci_list_head;  
    unsigned long nr_running;  
};
```

```
struct rq {  
    /* runqueue lock: */  
    raw_spinlock_t lock;  
  
    struct cfs_rq cfs;  
    struct rt_rq rt;  
    struct dl_rq dl;  
  
    struct domaci_rq dq;
```


domaci scheduler class

- Opisan u fajlu `/kernel/sched/domaci.c`
- Bitne funkcije:
 - `enqueue_task_domaci()`
 - `dequeue_task_domaci()`
 - `pick_next_task_domaci()`
 - `task_tick_domaci()`
 - ...

```
const struct sched_class domaci_sched_class = {  
    .next      = &fair_sched_class,  
    .enqueue_task    = enqueue_task_domaci,  
    .dequeue_task    = dequeue_task_domaci,  
    .yield_task     = yield_task_domaci,  
    .yield_to_task   = yield_to_task_domaci,  
  
    .check_preempt_curr = check_preempt_domaci,  
  
    .pick_next_task   = pick_next_task_domaci,  
    .put_prev_task    = put_prev_task_domaci,  
  
#ifdef CONFIG_SMP  
    .select_task_rq   = select_task_rq_domaci,  
    .migrate_task_rq  = NULL,  
  
    .rq_online        = NULL,  
    .rq_offline       = NULL,  
  
    .task_waking      = NULL,  
    .task_dead        = NULL,  
    .set_cpus_allowed = set_cpus_allowed_common,  
#endif
```

```
    .set_curr_task    = set_curr_task_domaci,  
    .task_tick        = task_tick_domaci,  
    .task_fork        = NULL,  
  
    .prio_changed     = prio_changed_domaci,  
    .switched_from    = NULL,  
    .switched_to      = switched_to_domaci,  
  
    .get_rr_interval  = get_rr_interval_domaci,  
  
    .update_curr      = update_curr_domaci,  
  
#ifdef CONFIG_FAIR_GROUP_SCHED  
    .task_move_group  = NULL,  
#endif  
};
```

domaci scheduler

- enqueue_task_domaci

```
static void
enqueue_task_domaci(struct rq *rq, struct task_struct *p, int flags)
{
    struct list_head *head_node = &rq->dq.domaci_list_head;

    if(!p->domaci_rt)
    {
        p->domaci_rt = jiffies;
        p->domaci_ticks = 0;
    }

    if(head_node->next != head_node)
    {
        p->domaci_list.prev = head_node->prev;
        p->domaci_list.next = head_node;
        head_node->prev->next = &p->domaci_list;
        head_node->prev = &p->domaci_list;
    }
    else
    {
        head_node->next = &p->domaci_list;
        head_node->prev = &p->domaci_list;
        p->domaci_list.next = head_node;
        p->domaci_list.prev = head_node;
    }

    rq->dq.nr_running++;
    add_nr_running(rq, 1);
}
```

- dequeue_task_domaci

```
static void dequeue_task_domaci(struct rq *rq, struct task_struct *p, int flags)
{
    p->domaci_list.prev->next = p->domaci_list.next;
    p->domaci_list.next->prev = p->domaci_list.prev;

    p->domaci_list.next = NULL;
    p->domaci_list.prev = NULL;

    rq->dq.nr_running--;
    sub_nr_running(rq, 1);
}
```

domaci scheduler

- pick_next_task_domaci

```
static struct task_struct *
pick_next_task_domaci(struct rq *rq, struct task_struct *prev)
{
    struct task_struct *task;
    struct list_head *head_node = &rq->dq.domaci_list_head;
    struct list_head *node;

    if(head_node->next == head_node)
        return NULL;

    task = list_entry(head_node->next, struct task_struct, domaci_list);

    node = head_node->next;
    dequeue_node_domaci(rq, node);
    enqueue_node_domaci(rq, node);

    task->domaci_time_slice = DOMACI_SLICE;
    return task;
}
```

- task_tick_domaci

```
static void task_tick_domaci(struct rq *rq, struct task_struct *p, int queued)
{
    p->domaci_time_slice--;
    p->domaci_ticks++;

    if(p->domaci_time_slice < 1)
        set_tsk_need_resched(p);
}
```

Sistemiški poziv

- Sistemiški poziv *sys_domaci* dodat u *domaci* scheduler
- Prikazivanje statistike procesa koji su se izvršili na *domaci* sched
- Statistika obuhvata:
 - pid procesa
 - vreme izvršenja na CPU
 - vreme provedeno od kreiranja do završenja procesa
- Funkcije:
 - *sys_domaci* - sistemiški poziv
 - *domaci_update_statistics* - računanje statistike

Sistemiški poziv

- sys_domaci

```
asm linkage long sys_domaci(unsigned long* niz)
{
    int i, ret = 0;
    unsigned long podaci[3 * PROCESS_BUFF_LEN] = {0};

    for(i = 0; i < PROCESS_BUFF_LEN; i++)
    {
        if(domaci_pids[i] > 0)
        {
            podaci[i] = domaci_pids[i];
            podaci[i + PROCESS_BUFF_LEN] = domaci_ticks[i];
            podaci[i + 2*PROCESS_BUFF_LEN] = domaci_rt[i];
        }
        else
            break;
    }

    if(copy_to_user(niz, podaci, sizeof(long) * 3 * PROCESS_BUFF_LEN))
        ret = -EFAULT;

    return ret;
}
```

- domaci_update_statistics

```
DEFINE_SPINLOCK(domaci_lock);

unsigned long domaci_index = 0;
unsigned int domaci_pids[PROCESS_BUFF_LEN] = {0};
unsigned long domaci_ticks[PROCESS_BUFF_LEN] = {0};
unsigned long domaci_rt[PROCESS_BUFF_LEN] = {0};

void domaci_update_statistics(struct task_struct *tsk)
{
    spin_lock(&domaci_lock);
    domaci_pids[domaci_index] = tsk->pid;
    domaci_rt[domaci_index] = jiffies - tsk->domaci_rt;
    domaci_ticks[domaci_index] = tsk->domaci_ticks;

    domaci_index = (domaci_index + 1) % PROCESS_BUFF_LEN;
    spin_unlock(&domaci_lock);
}
```


Sistemiški poziv

- Dodavanje sistemskog poziva u tabelu sistemskih poziva
 - **400 common domaci sys_domaci** - 64 bit (syscall_64.tbl)
 - **400 i386 domaci sys_domaci** - 32 bit (syscall_32.tbl)
- Dodavanje sistemskog poziva **sys_domaci** u header fajl sistemskih poziva (/include/linux/syscalls.h)
 - `asmlinkage long sys_domaci(unsigned long* niz);`

asmlinkage - tag koji kompajler koristi da pokupi podatke sa steka, a ne iz registra.

Testiranje

- Operativni sistemi
 - Manjaro 15.12
 - Ubuntu 15.10
- Linux kernel
 - 4.4.1
(www.kernel.org)



Koraci prilikom build-ovanja

- Dodavanje *domaci.o* u odgovarajući Makefile (/kernel/sched/Makefile)
 - obj-y += idle_task.o fair.o **domaci.o** rt.o deadline.o stop_task.o
- Manjaro build:
 - \$ make defconfig O=~/.linux/linux-build/
 - \$ make -jN CC="ccache gcc" O=~/.linux/linux-build/ > /dev/null
- Instalacija:
 - \$ cd ../linux-build
 - \$ su root
 - # make modules_install
 - # cp -v arch/x86/boot/bzImage /boot/vmlinuz-4.4.1
 - # mkinitcpio -k 4.4.1
 - c /etc/mkinitcpio.conf
 - g /boot/initramfs-4.4.1.img
 - # update-grub
- Ubuntu build:
 - \$ make defconfig
 - \$ make -jN CC="ccache gcc" deb-pkg
- Instalacija:
 - # dpkg -i linux-image-4.4.1.deb
 - # dpkg -i linux-headers-4.4.1.deb

Testiranje

- Dodeliti procesu da se izvršava na *domaci* scheduleru, pomoću sistemskog poziva *setscheduler()*.
- Kreiranje 10 deteta procesa (pomoću *fork()*).
- Dodati svim procesima kompleksne zadatke.
- Pomoću sistemskog poziva *sys_domaci()*, dobiti statistiku i odštampati je.

Korisnički program – part 1

```
1  #include <stdio.h>
2  #include <unistd.h>
3  #include <stdlib.h>
4  #include <sched.h>
5
6  #define KIDS 10
7  #define WORKLOAD 100000
8
9  int main()
10 {
11     int i = 0, j = 0, k;
12     pid_t p = 0;
13     long niz[3 * PROCESS_BUFF_LEN] = {0};
14
15     struct sched_param param;
16
17     param.sched_priority = 0;
18
19     if (sched_setscheduler(0, SCHED_DOMACI, &param) != 0)
20     {
21         perror("sched_setscheduler() failed!");
22         exit(EXIT_FAILURE);
23     }
24
25     printf("PARENT pid= %d, sched= %d\n\n", getpid(), sched_getscheduler(0));
26
```

Korisnički program - part 2

```
27     for(i = 0; i < KIDS; ++i)
28     {
29         p = fork();
30
31         if(p > 0)
32             continue;
33         else if(p == -1)
34         {
35             perror("fork() failed!");
36             continue;
37         }
38         else
39         {
40             for(j = 0; j < WORKLOAD; ++j)
41             {
42                 k = rand() * rand();
43                 k = rand() / rand();
44                 k = rand() - rand();
45                 k = rand() + rand();
46             }
47
48             printf("CHILD pid= %d, ppid= %d, sched= %d\n", getpid(), getppid(), sched_getscheduler(0));
49             exit(0);
50         }
51     }
52 }
53
```

Korisnički program - part 3

```
54     for(i = 0; i < KIDS; ++i)
55         wait();
56
57     if (syscall(400, niz) != 0)
58     {
59         perror("sys_domaci() failed!");
60         exit(EXIT_FAILURE);
61     }
62
63     printf("\nSYSCALL:\n");
64     for(i = 0; i < PROCESS_BUFF_LEN; ++i)
65     {
66         if(!niz[i])
67             break;
68
69         printf("pid= %d, ticks= %ld, rt= %ld\n", niz[i], niz[i + PROCESS_BUFF_LEN], niz[i + 2*PROCESS_BUFF_LEN]);
70     }
71
72     exit(0);
73 }
```

Korisnički program

Rezultat izvršenja

```

.                                     sched
PARENT pid= 3174, sched= 10

CHILD pid= 3175, ppid= 3174, sched= 10
CHILD pid= 3176, ppid= 3174, sched= 10
CHILD pid= 3177, ppid= 3174, sched= 10
CHILD pid= 3178, ppid= 3174, sched= 10
CHILD pid= 3179, ppid= 3174, sched= 10
CHILD pid= 3180, ppid= 3174, sched= 10
CHILD pid= 3181, ppid= 3174, sched= 10
CHILD pid= 3182, ppid= 3174, sched= 10
CHILD pid= 3183, ppid= 3174, sched= 10
CHILD pid= 3184, ppid= 3174, sched= 10

SYSCALL:
pid= 3175, ticks= 18, rt= 153
pid= 3176, ticks= 18, rt= 156
pid= 3177, ticks= 18, rt= 159
pid= 3178, ticks= 17, rt= 161
pid= 3179, ticks= 18, rt= 164
pid= 3180, ticks= 18, rt= 167
pid= 3181, ticks= 17, rt= 169
pid= 3182, ticks= 18, rt= 172
pid= 3183, ticks= 17, rt= 174
pid= 3184, ticks= 18, rt= 177

Process returned 0 (0x0)   execution time : 0.184 s
Press ENTER to continue.

```

\$ dmesg

```
[ 66.299286] fja enqueue_task_domaci(3174)
[ 66.299378] fja enqueue_task_domaci(3175)
[ 66.299426] fja enqueue_task_domaci(3176)
[ 66.299471] fja enqueue_task_domaci(3177)
[ 66.299513] fja enqueue_task_domaci(3178)
[ 66.299559] fja enqueue_task_domaci(3179)
[ 66.299601] fja enqueue_task_domaci(3180)
[ 66.299649] fja enqueue_task_domaci(3181)
[ 66.299690] fja enqueue_task_domaci(3182)
[ 66.299736] fja enqueue_task_domaci(3183)
[ 66.299781] fja enqueue_task_domaci(3184)
[ 66.299802] fja dequeue_task_domaci(3174)
[ 66.299804] fja pick_next_task_domaci picked(3175)
[ 66.303994] fja pick_next_task_domaci picked(3176)
[ 66.308991] fja pick_next_task_domaci picked(3177)
[ 66.313991] fja pick_next_task_domaci picked(3178)
[ 66.318991] fja pick_next_task_domaci picked(3179)
[ 66.323991] fja pick_next_task_domaci picked(3180)
[ 66.328992] fja pick_next_task_domaci picked(3181)
[ 66.333992] fja pick_next_task_domaci picked(3182)
[ 66.338988] fja pick_next_task_domaci picked(3183)
[ 66.343988] fja pick_next_task_domaci picked(3184)
[ 66.349004] fja pick_next_task_domaci picked(3175)
[ 66.353993] fja pick_next_task_domaci picked(3176)
[ 66.358991] fja pick_next_task_domaci picked(3177)
[ 66.363990] fja pick_next_task_domaci picked(3178)
[ 66.368989] fja pick_next_task_domaci picked(3179)
[ 66.373989] fja pick_next_task_domaci picked(3180)
[ 66.378989] fja pick_next_task_domaci picked(3181)
[ 66.383989] fja pick_next_task_domaci picked(3182)
[ 66.388989] fja pick_next_task_domaci picked(3183)
[ 66.393990] fja pick_next_task_domaci picked(3184)
[ 66.398998] fja pick_next_task_domaci picked(3175)
[ 66.403990] fja pick_next_task_domaci picked(3176)
[ 66.408989] fja pick_next_task_domaci picked(3177)
[ 66.413990] fja pick_next_task_domaci picked(3178)
[ 66.418990] fja pick_next_task_domaci picked(3179)
[ 66.423990] fja pick_next_task_domaci picked(3180)
[ 66.428990] fja pick_next_task_domaci picked(3181)
[ 66.433991] fja pick_next_task_domaci picked(3182)
[ 66.438990] fja pick_next_task_domaci picked(3183)
[ 66.443993] fja pick_next_task_domaci picked(3184)
[ 66.448992] fja pick_next_task_domaci picked(3175)
[ 66.452646] pid=3175 rt=153 ticks=18
[ 66.452651] fja enqueue_task_domaci(3174)
[ 66.452657] fja dequeue_task_domaci(3175)
[ 66.452658] fja pick_next_task_domaci picked(3176)
[ 66.455361] pid=3176 rt=156 ticks=18
[ 66.455367] fja dequeue_task_domaci(3176)
[ 66.455369] fja pick_next_task_domaci picked(3177)
[ 66.458023] pid=3177 rt=159 ticks=18
[ 66.458028] fja dequeue_task_domaci(3177)
[ 66.458030] fja pick_next_task_domaci picked(3178)
[ 66.460663] pid=3178 rt=161 ticks=17
[ 66.460668] fja dequeue_task_domaci(3178)
[ 66.460669] fja pick_next_task_domaci picked(3179)
[ 66.463316] pid=3179 rt=164 ticks=18
[ 66.463321] fja dequeue_task_domaci(3179)
[ 66.463323] fja pick_next_task_domaci picked(3180)
[ 66.466025] pid=3180 rt=167 ticks=18
[ 66.466033] fja dequeue_task_domaci(3180)
[ 66.466034] fja pick_next_task_domaci picked(3181)
[ 66.468668] pid=3181 rt=169 ticks=17
[ 66.468674] fja dequeue_task_domaci(3181)
[ 66.468675] fja pick_next_task_domaci picked(3182)
[ 66.471273] pid=3182 rt=172 ticks=18
[ 66.471279] fja dequeue_task_domaci(3182)
[ 66.471280] fja pick_next_task_domaci picked(3183)
[ 66.473862] pid=3183 rt=174 ticks=17
[ 66.473867] fja dequeue_task_domaci(3183)
[ 66.473868] fja pick_next_task_domaci picked(3184)
[ 66.476582] pid=3184 rt=177 ticks=18
[ 66.476588] fja dequeue_task_domaci(3184)
[ 66.476589] fja pick_next_task_domaci picked(3174)
[ 66.476739] pid=3174 rt=177 ticks=0
[ 66.476748] fja dequeue_task_domaci(3174)
```